

PRD: View Assessments on Admin Side

1. Overview

The assessment reporting surface within the admin course view is due for a ground-up redesign. The current experience — a static table with emoji icons, inconsistent modals, and no summary statistics — has been described as outdated and falls short of the expectations of L&D / People admins who periodically review cohort performance. This ticket replaces that experience with a modern, insight-first reporting interface that treats assessments as a reporting surface rather than a data dump. The redesign introduces AI-generated insights, contextual summary statistics, a dual-pivot view (by assessment / by learner), and a side drawer for sequential detail review — all within the existing Assessments tab rather than a modal.

Grading, commenting, editing, and reassigning are explicitly out of scope (V1 is read-only plus CSV export). Admin scoring for short-text and exercise types is deferred to V2, and the V1 data model deliberately excludes correctness fields for those types so scoring cannot leak into the UI.

2. Jobs To Be Done

Main Job

When a cohort finishes a course, the L&D admin needs to understand how learners performed across all assessments so they can identify knowledge gaps, report on program effectiveness, and decide whether to intervene.

Job Map

Stage	What the admin does	Where they struggle today
Define	Decide which course and cohort to evaluate	No issues — course selection exists
Locate	Navigate to the Assessments tab	Tab exists but offers no summary; the admin has no signal until they click into each row
Prepare	Scan for overall performance signals	No summary statistics at all; the Correct % column shows em dashes for 3 of 5 visible types, making scanning unreliable
Confirm	Drill into a specific assessment to validate a concern	Opens a 720 px modal that is a dead end — no previous/next, so the admin must close and re-open for every assessment
Execute	Review individual responses, compare across learners	Modal anatomy differs per type (stat cards + pagination for MCQ, nothing for exercise); no "by learner" view exists

Monitor	Track patterns over time, share with stakeholders	No AI insights, no exportable summary; the only export is a "Download responses" button styled as a primary CTA because the modal template forces it
Resolve	Take action — adjust content, follow up with learners	No actionable insight layer; the admin must mentally synthesise raw data

Related Jobs

- **Report to leadership on L&D program ROI** — the admin needs exportable, digestible summaries, not raw CSVs alone.
- **Identify at-risk learners early** — a "by learner" pivot lets the admin spot individuals who struggled across multiple assessments.
- **Improve course content iteratively** — AI insights that surface where learners struggled inform content revisions.

Emotional & Social Dimensions

- **Emotional:** The admin wants to feel *confident and in control* — that they have a complete picture of cohort performance without manually assembling it. The current em-dash-riddled table and dead-end modals create anxiety and distrust.
- **Social:** The admin wants to be perceived as *data-driven and thorough* by leadership and course designers. Being able to surface AI-generated insights elevates their role from "person who downloads CSVs" to "learning strategist."

3. Goals

1. **Insight-first experience:** An admin should grasp the headline performance story within five seconds of landing on the Assessments tab, without clicking into any individual assessment.
2. **Type-appropriate metrics:** Every assessment type displays a meaningful result metric (percent correct, vote distribution, response/file count) — no em dashes, ever.
3. **Seamless sequential review:** An admin can drill into an assessment and navigate previous/next through the full list without returning to the table.
4. **Dual-pivot analysis:** The admin can switch between "by assessment" and "by learner" views to answer different questions from the same data.
5. **AI-powered synthesis:** A single-click AI Insights card surfaces where learners struggled and what they mastered, with a clear generate → generating → generated lifecycle.
6. **Clean export:** CSV export of all answers is available and discoverable without dominating the UI.
7. **Scalable design:** The interface handles tens of assessments and hundreds of responses gracefully, with search, type filtering, and pagination.

8. **Additive type support:** The shared chrome + type-specific body architecture means a tenth assessment type can be added without a new template.

4. Job Stories

Orientation & Summary

- **When** I open the Assessments tab after a cohort has completed a course, **I want to** see summary statistics (response count vs. enrolment, average score, completion rate) at a glance, **so I can** immediately gauge overall cohort performance without drilling into individual assessments.
- **When** I see the summary stat cards, **I want** each card to show the metric appropriate to the assessment type (percent correct for auto-graded, vote distribution for polls, response count for short text/exercise), **so I can** trust the data and never encounter meaningless placeholders.

AI Insights

- **When** I need to report on where learners struggled and what they mastered, **I want to** generate an AI-powered insight summary with one click, **so I can** get a synthesised narrative instead of manually reviewing every response.
- **When** the AI insight is generating, **I want to** see progress context (e.g., "Analysing 47 responses...") and skeleton placeholders, **so I can** understand the system is working and estimate how long it will take.
- **When** the AI insight has been generated, **I want to** see the timestamp of generation, **so I can** judge whether it reflects the latest submissions or needs refreshing.

Detail Review

- **When** I want to review a specific assessment's responses, **I want to** open a side drawer that shows the detail without leaving the table, **so I can** keep the list context visible and reference other assessments.
- **When** I am reviewing an assessment in the drawer, **I want to** navigate to the previous or next assessment using arrows and see my position (e.g., "3 of 12"), **so I can** work through all assessments sequentially without closing and reopening.
- **When** I am reviewing responses in the drawer, **I want to** see learner identity and role alongside each response, **so I can** follow up with specific individuals if needed.

Pivot & Filter

- **When** I want to understand a specific learner's performance across all assessments, **I want to** switch to a "By learner" view, **so I can** see their complete assessment profile in one place.

- **When** the course has many assessments, **I want to** search by title, filter by type, and paginate, **so I can** quickly find the assessment I care about.

Export

- **When** I need to share raw response data with a stakeholder or import it into another tool, **I want to** export all answers as a CSV, **so I can** get the data out without screenshots or manual copying.

5. Requirements

Functional Requirements

FR-1: Summary Stat Cards

- Display a row of stat cards at the top of the Assessments tab summarising cohort-level performance.
- Minimum cards: total assessments, overall response count vs. enrolment denominator, average score (auto-graded types only), completion rate.
- Each card must show the type-appropriate metric; if a metric is not applicable to a type, that type is excluded from the card's calculation rather than showing an em dash.

FR-2: AI Insights Card

- A dedicated card below the stat row with an explicit three-state lifecycle:
 - **Generate:** A "Generate Insights" button when no insight exists or the admin wants to refresh.
 - **Generating:** Skeleton placeholder with contextual message (e.g., "Analysing 47 responses..."). No bare spinner.
 - **Generated:** The insight text with a "Generated at [timestamp]" label.
- Insights should surface: where learners struggled, what they mastered, and any notable patterns.
- The insight must be clearly labelled as AI-generated.

FR-3: Assessment List (By Assessment View)

- A filterable, searchable, paginated data table replacing the current table.
- Columns: Title, Type (with design-system-compliant icons, not emoji), Responses (as "n / enrolled"), Result (type-appropriate: % correct, vote distribution, response/file count).
- Search by assessment title.
- Filter by assessment type (multi-select).
- Pagination for lists exceeding a configurable page size.
- Row click opens the side drawer.

FR-4: By Learner View

- A toggle/segmented control switches between "By assessment" and "By learner" pivots.
- The By Learner view lists learners with their aggregated performance across all assessments.
- Row click opens a side drawer showing that learner's responses per assessment.

FR-5: Side Drawer (Detail View)

- Opens from the right, overlaying the table (overlay variant per the design system's Side Drawer at 720 px).
- Contains:
 - Assessment title and type badge.
 - Type-specific result display (e.g., bar chart for MCQ distribution, list of text responses for short text).
 - Individual responses with learner name and role.
 - Submission timestamps and attempt history (demoted to drawer per stakeholder direction).
- Previous/Next navigation arrows with "n of m" position indicator.
- Keyboard shortcut support (arrow keys) for sequential navigation.
- Close button and click-outside-to-close behaviour.

FR-6: CSV Export

- A secondary action (e.g., icon button or dropdown item) on the Assessments tab — not a primary CTA.
- Exports all responses across all assessments for the course in a single CSV.
- Includes: assessment title, type, learner name, learner role, response, result (where applicable).

FR-7: Shared Chrome, Type-Specific Bodies

- The side drawer chrome (header, navigation, close, export) is shared across all nine assessment types.
- The body area renders a type-specific component:
 - **Multiple Choice:** Option distribution bar/chart, correct answer highlighted, per-response list.
 - **Match the Pairs:** Pair accuracy summary, per-response correctness.
 - **Sequence:** Sequence accuracy summary, per-response correctness.
 - **Categorise:** Category accuracy summary, per-response correctness.
 - **Fill in the Blanks:** Blank-level accuracy, per-response list.
 - **Short Text:** Response list only (no correctness field in V1).
 - **Exercise (file upload):** File list with download links (no correctness field in V1).
 - **Poll:** Vote distribution chart, per-response list.

- **Situational Test:** Option distribution, scoring summary, per-response list.
- Adding a tenth type requires only a new body component, no changes to chrome.

FR-8: Empty & Edge States

- Zero responses: stat cards show "0 / [enrolled]", AI Insights button is disabled with tooltip "No responses yet."
- Single response: all views functional; AI Insights available but may note limited data.
- No assessments in course: empty state with guidance message.

FR-9: No Correctness for Short Text & Exercise

- The UI must not display, imply, or reserve space for a correctness metric on Short Text or Exercise types.
- This aligns with the V1 data model where these types have no correctness field, preventing any accidental scoring leakage before V2.

6. Research & Best Practices

Synthesis

The redesign is grounded in a clear industry consensus: admin reporting surfaces should lead with insight, not raw data. The UXPin dashboard design guide establishes the **five-second rule** — users should grasp the main insight within five seconds — and recommends limiting visible KPIs to 3–5 items to respect working memory limits, deferring secondary metrics via progressive disclosure ([UXPin](#)). This directly informs the stat card row: a small number of high-signal cards (response rate, average score, completion) above the table, with detailed breakdowns deferred to the drawer.

The **AI Insights card** draws on multiple precedents. Google Forms now offers a "Summarise with Gemini" button that generates a concise summary of key insights across all responses with one click — directly analogous to our generate → generating → generated lifecycle ([xFanatical](#)). AufaitUX's enterprise AI patterns emphasise **Explainability Layers** and **Human-Verified vs AI-Generated labeling**, both of which mandate that our AI insight card is clearly marked as AI-generated with a timestamp ([AufaitUX](#)). Eleken's SaaS AI dashboard guide notes that embedded workflow modules (insight → action) cut response time by 42% per Forrester 2025 ([Eleken](#)). GitNexa's SaaS patterns survey confirms the trend: Gartner predicts 60% of SaaS analytics platforms will integrate generative AI interfaces by 2027 ([GitNexa](#)).

The decision to **replace modals with a side drawer** is strongly supported by NN/g's data tables research, which explicitly states that modals "cover adjacent records and the user won't be able to reference or copy data from a similar record," favouring non-modal side panels for detail views ([NN/g Data Tables](#)). NN/g's progressive disclosure article reinforces this: the master-detail side panel is itself a form of progressive disclosure, with the table as summary and the panel as detail ([NN/g Progressive Disclosure](#)). The UX Planet case study from Grubtech found that replacing modals with layered drawers improved context preservation, concluding that "modals don't scale when depth, flexibility, and memory are

required" ([UX Planet](#)). PatternFly's drawer guidelines confirm the overlay variant (floats over content) as appropriate for primary-detail layouts spanning the height of primary content ([PatternFly](#)).

The **previous/next navigation** pattern is well-established. DataTables Editor demonstrates previous/next buttons that save and load sequentially without closing the detail view ([DataTables Editor](#)). Oracle Sales Cloud implements next/previous arrows with a "list name" return link when drilling into list items ([Oracle](#)). The "n of m" indicator is critical for orientation.

The **by assessment vs. by learner pivot** is standard in LMS assessment reporting. Educate-me's LMS reporting guide identifies this dual pivot as standard across Blackboard, LearnDash, and Instancy, alongside key metrics: completion rate, pass/fail ratio, average score, and question-level discrimination index ([Educate-me](#)). InsightsRoom demonstrates the market shift from raw data tables to insight-first reporting with auto-generated dashboards and interactive filtering ([InsightsRoom](#)). The admin dashboard best practices guide on Medium reinforces that best-in-class dashboards prioritize clarity, conciseness (5–7 core KPIs), and contextual insights that tell a story ([Medium](#)).

SaaSFrame catalogues 62 side panel examples from SaaS products, confirming the pattern maximises screen real estate by keeping users in context ([SaaSFrame](#)). The competitive landscape, including Sana Labs and WorkRamp, confirms that AI-driven analytics dashboards are table stakes for modern L&D platforms ([Teachfloor](#)).

UX References

App	Flow / Screen	URL	Pattern Description	Relevance
Deel	Admin dashboard with stat cards and data table	View	Summary stat cards above a filterable data table	Direct precedent for the stat cards + assessment list layout on the Assessments tab
Vanta	Dashboard with summary statistics and table	View	Compliance/security dashboard with stats and list	Mentioned in the ticket's competitor analysis; validates the stat-card-above-table pattern
Uxcel	Assessment/learning dashboard	View	Learning platform admin with results display	LMS-adjacent product showing assessment-like reporting surfaces
Maze	Side drawer detail panel	View	Drawer for research/test result details	Mentioned in the ticket's competitor analysis; validates drawer-based detail for research/test results
Workable	Side drawer for record detail	View	HR/admin side panel for candidate/record details	Closest persona match — HR admin reviewing records via side drawer

Asana	Detail panel alongside list	View	Master-detail with persistent list context	Well-known side panel pattern preserving list context during detail review
Fibery	Side panel for list item detail	View	Inline detail panel with main view visible	Shows how a detail panel can push content rather than overlay
Sprig	AI insights in analytics dashboard	View	AI-generated analysis cards within reporting UI	Direct precedent for the AI Insights card with generate/display lifecycle
Dovetail	AI insights in research analytics	View	AI-powered insight summaries for qualitative data	Shows AI analysis of open-ended responses — analogous to short-text/exercise assessment types
Amplitude	Analytics dashboard with generated insights	View	AI-assisted analytics with data visualisations	Enterprise analytics showing how AI insights sit alongside charts and tables
Productboard	AI-generated insights in reporting	View	AI insights card within product analytics view	Shows the generate → generated lifecycle for AI analysis
Hotjar	Tab/view switcher in analytics	View	Segmented control switching between data views	Direct precedent for the "By assessment" vs "By learner" switcher
Google Analytics	Tab navigation between data views	View	Tab bar toggling between analytics perspectives	Standard pattern for switching data pivots in a reporting interface
Gamma	View switcher in reporting/analytics	View	Toggle between presentation views of same data	Shows how to switch between item-centric vs audience-centric views

Built for Mars Lessons

Title	Type	URL	Lesson	Relevance
-------	------	-----	--------	-----------

Post-video analytics (Loom)	UX Bite	View	Loom surfaces motivating analytics in visually distinct card-based modules — presenting derived insights in card form, not just raw data, makes analytics feel actionable and engaging.	Direct pattern for the AI Insights card: surface derived insights (where learners struggled, what they mastered) as standalone cards above the data table, not buried in it.
Global summary settings (Zillow)	UX Bite	View	Zillow lets users customise which summary facts appear on property cards, making browsing feel more efficient and specific to their needs.	Applicable to the stat card row: let the type-appropriate metric surface contextually rather than forcing a single metric template that results in em-dash placeholders for incompatible types.
The UX Psychology of Waiting (and Loading)	Cheatsheet	View	Effective loading states add variability — display rotating tips, relevant context, or preview skeletons rather than bare spinners. Meaningful loading content reduces perceived wait time.	Directly applicable to the AI Insights card's generate → generating → generated lifecycle: show skeleton cards, progress context ("Analysing 47 responses..."), and avoid jarring layout shifts when insights arrive.
BFM Synthesised Advice	Analysis	N/A	Three key recommendations: (1) Use modular insight cards above data tables, styled distinctly from raw data — inspired by Loom; (2) For AI generation loading states, show skeleton placeholders with contextual messages rather than spinners, and display "n of m" position plus keyboard nav in drawers; (3) Allow users to customise which summary stats surface, inspired by Zillow.	Synthesised guidance covering all three core UX challenges: AI insight presentation, loading lifecycle, and drawer navigation.

7. Plan of Action

Phase 1: Foundation — Stat Cards & Modernised Table

- ☐ Audit the existing Assessments tab component tree and data-fetching layer; document current API contracts.
- ☐ Design and implement summary stat cards (total assessments, response rate vs. enrolment, average score, completion rate) with type-aware metric logic (no em dashes).
- ☐ Replace emoji type icons with design-system-compliant icons across all nine types.
- ☐ Implement search-by-title, multi-select type filter, and pagination on the assessment list table.
- ☐ Ensure the Correct % column is replaced with a "Result" column showing type-appropriate metrics.
- ☐ Implement empty and edge states (zero responses, no assessments).

Phase 2: Side Drawer & Sequential Navigation

- ☐ Build the shared drawer chrome: header (title + type badge), close button, previous/next arrows with "n of m" indicator.
- ☐ Implement keyboard navigation (arrow keys for previous/next, Escape to close).
- ☐ Build type-specific body components for all nine assessment types with individual response lists (learner name + role).
- ☐ Integrate submission timestamps and attempt history within the drawer.
- ☐ Ensure click-outside-to-close and maintain list scroll position on drawer open/close.
- ☐ Ensure Short Text and Exercise body components render no correctness metric per FR-9.

Phase 3: By Learner View & Pivot Switcher

- ☐ Design and implement the segmented control ("By assessment" / "By learner") at the top of the tab.
- ☐ Build the By Learner list view with aggregated per-learner performance.
- ☐ Wire the By Learner row click to open the drawer with that learner's assessment-by-assessment responses.
- ☐ Ensure search, filter, and pagination work in both pivot views.

Phase 4: AI Insights Card

- ☐ Design the AI Insights card with three-state lifecycle UI: generate button → skeleton with contextual message → generated insight with timestamp.
- ☐ Build the backend endpoint to generate insights from assessment responses (integration with AI service).
- ☐ Implement "AI-generated" labelling on the card per enterprise AI UX patterns.

- ☐ Handle edge cases: disable generation when zero responses; show "limited data" caveat for very few responses.
- ☐ Implement refresh/regenerate capability.

Phase 5: CSV Export & Polish

- ☐ Implement CSV export as a secondary action (icon button or toolbar dropdown) — not a primary CTA.
- ☐ Include all required fields: assessment title, type, learner name, role, response, result.
- ☐ Conduct cross-browser and responsive testing.
- ☐ Conduct accessibility audit (keyboard navigation, screen reader support, ARIA labels on drawer and stat cards).
- ☐ Conduct usability testing with L&D admins reviewing a course with 15+ assessments and 50+ learners.

8. Risks & Mitigations

Risk	Impact	Mitigation
AI insight quality is poor or misleading for small cohorts	Admin loses trust in the feature; reputational risk	Disable generation below a minimum response threshold; add "limited data" caveat for small samples; clearly label as AI-generated with timestamp
Nine type-specific body components create high implementation cost	Phase 2 timeline overruns	Prioritise the four most-used types first (Multiple Choice, Short Text, Exercise, Poll); ship remaining five in a fast-follow; shared chrome ensures consistency regardless of body readiness
Side drawer at 720 px may crowd the table on smaller screens	Poor experience on laptops with ≤ 1280 px viewports	Use overlay (not inline push) variant so the table is not compressed; test at 1280 px breakpoint; consider responsive collapse to near-full-width drawer on small screens
"By learner" view requires a new API aggregation endpoint	Backend dependency delays frontend work	Stub the frontend with mock data behind a feature flag; define the API contract in Phase 1 so backend can work in parallel
Previous/next navigation desyncs if the underlying list changes (new responses arrive)	Admin sees stale data or skips an item	Snapshot the list order when the drawer opens; show a subtle "List updated — refresh?" indicator if the underlying data changes
CSV export for large courses (hundreds of	Admin experiences timeout or assumes	Generate CSV asynchronously; show progress indicator; consider server-side generation with a

learners x tens of assessments) is slow	failure	download link delivered via notification
V2 scoring for Short Text / Exercise requires data model changes	V1 architecture doesn't accommodate V2	V1 deliberately omits correctness fields for these types (unrepresentable, not hidden); V2 migration adds the field, and the body component conditionally renders scoring UI — no V1 refactoring needed
Design system lacks a documented full-page or large-panel component	Inconsistency with existing patterns	Use the documented Side Drawer (720 px) from overlays.md; avoid inventing undocumented components; propose a design system contribution if a wider variant is needed later